

Computer Architecture and Programming

ICS312 - Spring 2014
Machine-level and
Systems Programming

Henri Casanova (henric@hawaii.edu)


“Computer Architecture”?

- The field of Computer Architecture is about the fundamental structure of computer systems
 - What are the components
 - How are they interconnected?
 - How fast does the whole system operate?
 - How much power does it consume?
 - How much does it cost to mass-produce?
 - How to achieve desired speed/power/cost trade-offs?
- The conceptual model for computer architecture, that hasn't fundamentally changed since 1965: the **Von-Neumann architecture**


Von-Neumann

- In 1944, John von Neumann joined ENIAC
- He wrote a memo about computer architecture, formalizing ENIAC ideas
 - Eckert and Mauchly have pretty much been forgotten (they were in the trenches)
- These ideas became the **Von Neumann architecture model**
 - A **processor** that performs operations and controls all that happens
 - A **memory** that contains **code and data**
 - **I/O** of some kind

Von-Neumann Model



- Amazingly, it's still possible to think of the computer this way at a conceptual level (model from ~70 years ago!!!)
- But a computer today doesn't look quite like this



Data Stored in Memory

- All “information” in the computer is in binary form
 - Since Claude Shannon’s M.S. thesis in the 30’s
 - 0: zero voltage, 1: positive voltage (e.g., 5V)
 - bit: the smallest unit of information (0 or 1)
- The basic unit of memory is a byte
 - 1 Byte = 8 bits, e.g., “0101 1101”
 - 1 KiB = 2^{10} byte = 1,024 bytes
 - 1 MiB = 2^{10} KiB = 2^{20} bytes (~ 1 Million)
 - 1 GiB = 2^{10} MiB = 2^{30} bytes (~ 1 Billion)
 - 1 TiB = 2^{10} GiB = 2^{40} bytes (~ 1 Trillion)
 - 1 PiB = 2^{10} TiB = 2^{50} bytes (~ 1000 Trillion)
 - 1 EiB = 2^{10} PiB = 2^{60} bytes (~ 1 Million Trillion)
 - ...
- Note the “i” in the notations above: means “power of 2”

Data Stored in Memory

- Each byte in memory is labeled by a unique **address**
- An address is a number which identifies the memory location of each byte in memory
 - e.g., the byte at address 3 is 00010010
 - e.g., the byte at address 241 is 10110101
- Typically, we write address in binary as well
 - e.g., the byte at address 00000011 is 00010010
 - e.g., the byte at address 11110001 is 10110101
- All addresses in the machine have the same number of bits
 - e.g., 8-bit addresses
- The processor has instructions that say “Read the byte at address X and give me its value” and “Write some value into the byte at address X”

Conceptual View of Memory

address	content
0000 0000 0000 0000	0110 1110
0000 0000 0000 0001	1111 0100
0000 0000 0000 0010	0000 0000
0000 0000 0000 0011	0000 0000
0000 0000 0000 0100	0101 1110
0000 0000 0000 0101	1010 1101
0000 0000 0000 0110	0000 0001
0000 0000 0000 0111	0100 0000
0000 0000 0000 1000	1111 0101
...	...

Conceptual View of Memory

address	content
0000 0000 0000 0000	0110 1110
0000 0000 0000 0001	1111 0100
0000 0000 0000 0010	0000 0000
0000 0000 0000 0011	0000 0000
0000 0000 0000 0100	0101 1110
0000 0000 0000 0101	
0000 0000 0000 0110	
0000 0000 0000 0111	
0000 0000 0000 1000	1111 0101
...	...

At address 0000 0000 0000 0010 the content is 0000 0000

Conceptual View of Memory

address	content
0000 0000 0000 0000	0110 1110
0000 0000 0000 0001	1111 0100
0000 0000 0000 0010	0000 0000
0000 0000 0000 0011	0000 0000
0000 0000 0000 0100	0101 1110
0000 0000 0000 0101	
0000 0000 0000 0110	
0000 0000 0000 0111	
0000 0000 0000 1000	1111 0101
...	...

**At address 0000 0000 0000 0100
the content is 0101 1110**

Both Code and Data in Memory

- Once a program is loaded in memory, its address space contains both **code** and **data**
- To the CPU those are not really different, but the programmer knows which bytes are data and which are code
 - Always conveniently hidden from you if you've never written assembly
 - But we'll have to keep code/data straight in these lecture notes

Example Address Space

Address	Value
0000 1100	0110 1011
0000 1101	1111 0010
0000 1110	0010 0001
...	...
1000 0000	1111 0000
...	...
1111 0010	0101 1111


Code (addresses 0000 1100 to 0000 1110)
Data (addresses 1000 0000 to 1111 0010)

Memory


We need a CPU

- So now we have a memory in which we can store/retrieve bytes at precise location
- These bytes presumably have some useful meaning to us
 - e.g., integers, ASCII codes of characters, floating point numbers, RGB values
 - e.g., instructions that specify what to do with the data; when you buy a processor, the vendor defines the instruction set (e.g., instruction "0010 1101" means "increment some useful counter")
- The CPU is the piece of hardware that modifies the content of memory
 - In fact, one can really think of the CPU as a device that takes use from on memory state (i.e., all the stored content) to another memory state (some new, desired stored content)


What's in the CPU?



What's in the CPU?




What's in the CPU?



Registers: the "variables" that hardware instructions work with

Data can be loaded from memory into a register
 Data can be stored from a register back into memory
 Operands and results of computations are in registers
 Accessing a register is really fast
 There is a limited number of registers

What's in the CPU?




Arithmetic and Logic Unit: what you do computation with

used to compute a value based on current register values and store the result back into a register

+, *, /, -, OR, AND, XOR, etc.


What's in the CPU?



Program Counter: Points to the next instruction


Special register that contains the address in memory of the next instruction that should be executed
 (gets incremented after each instruction, or can be set to whatever value whenever there is a change of control flow)

What's in the CPU?



Current Instruction: Holds the instruction that's currently being executed

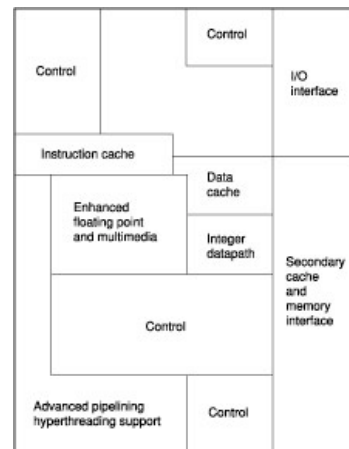
What's in the CPU?



Control Unit: Decodes instructions and make them happen

Logic hardware that decodes instructions (i.e., based on their bits) and sends the appropriate (electrical) signals to hardware components in the CPU

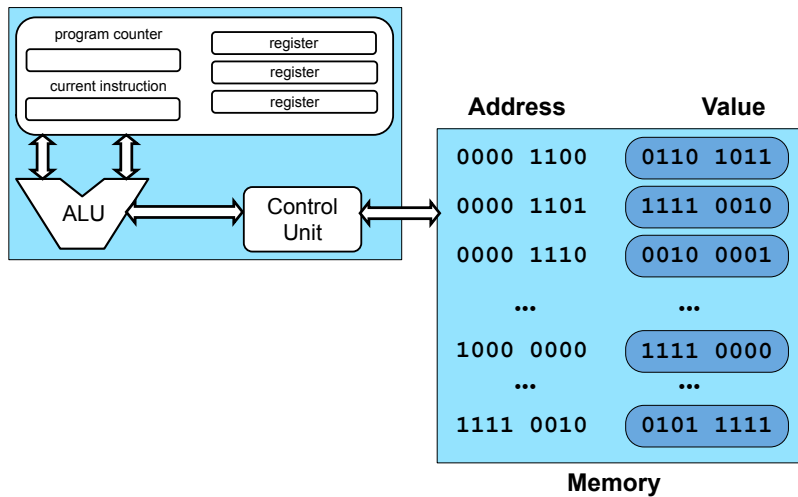
The CPU in its "Glory"



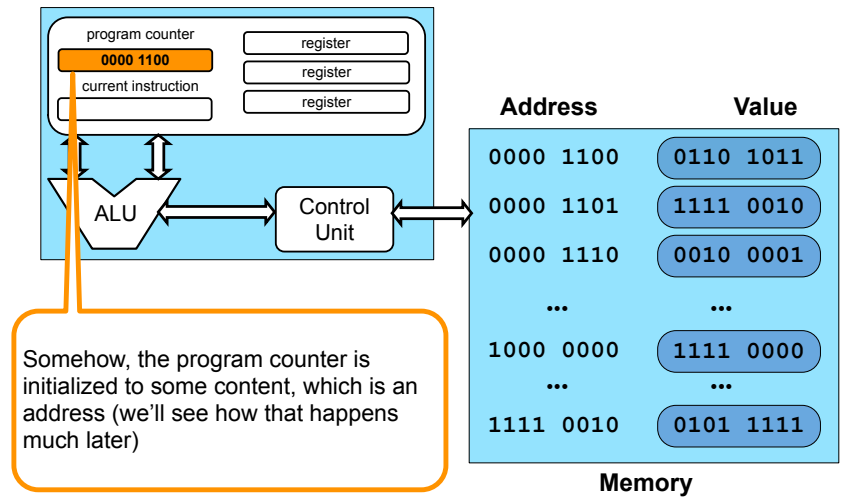
Fetch-Decode-Execute Cycle

- The **Fetch-Decode-Execute** cycle
 - The **control unit** fetches the next program instruction from memory
 - Using the **program counter** to figure out where that instruction is located in the memory
 - The instruction is decoded and signals are send to hardware components
 - e.g., is the instruction loading something from memory? is it adding two register values together?
 - Operands are fetched from **memory** and put in **registers**, if needed
 - The ALU executes computation, if any, and store results in the **registers**
 - Register values are stored back to **memory**, if needed
 - Repeat
- Computers today implement MANY variations on this model
- But one can still program with the above model in mind
 - but certainly without (fully) understanding performance issues

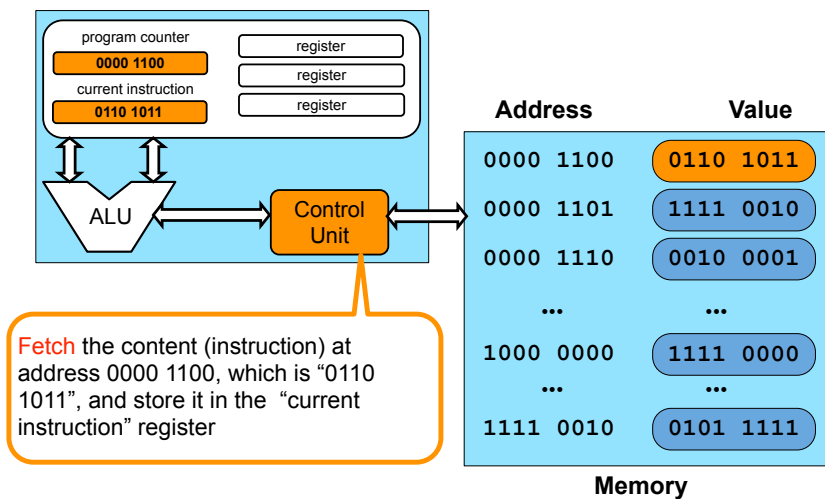
Fetch-Decode-Execute



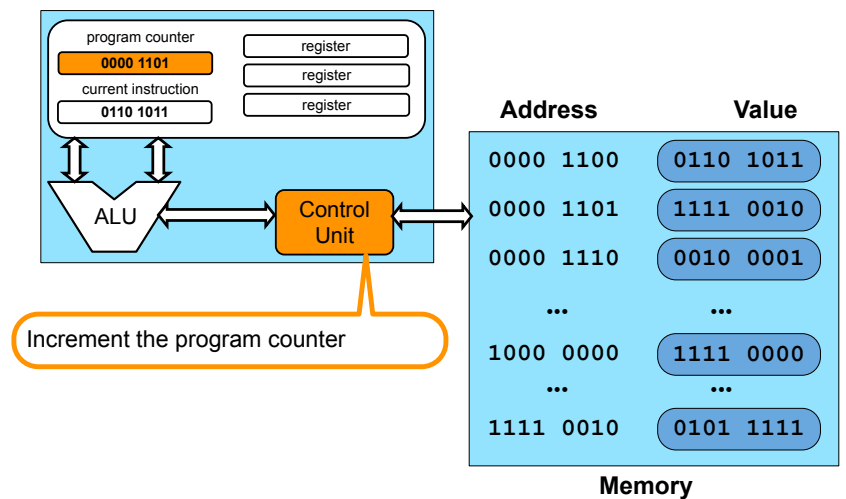
Fetch-Decode-Execute



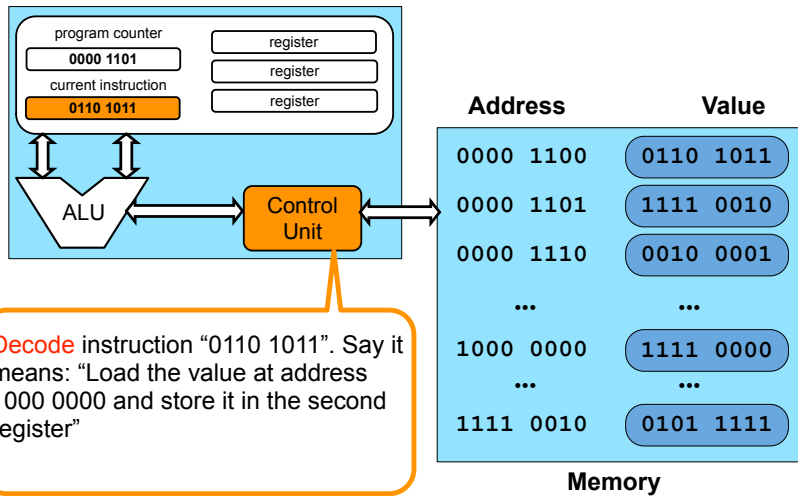
Fetch-Decode-Execute



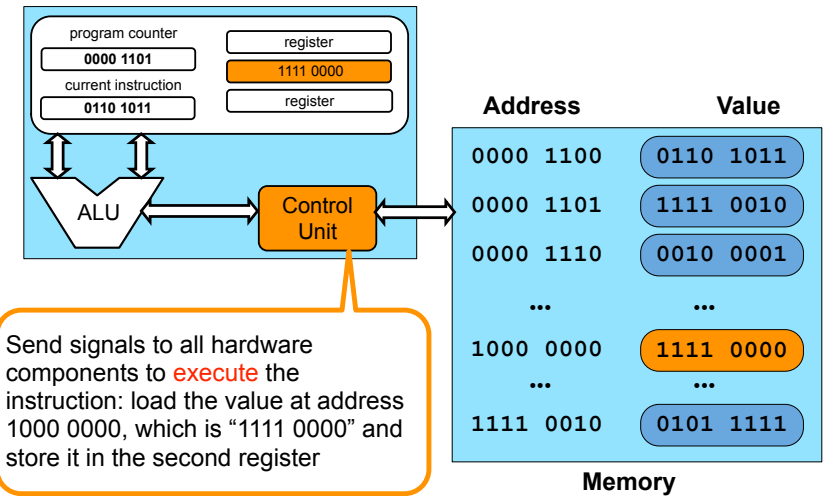
Fetch-Decode-Execute



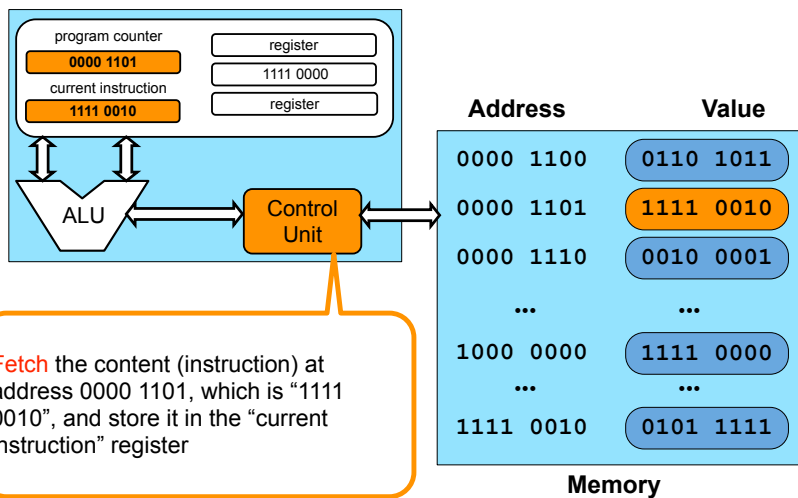
Fetch-Decode-Execute



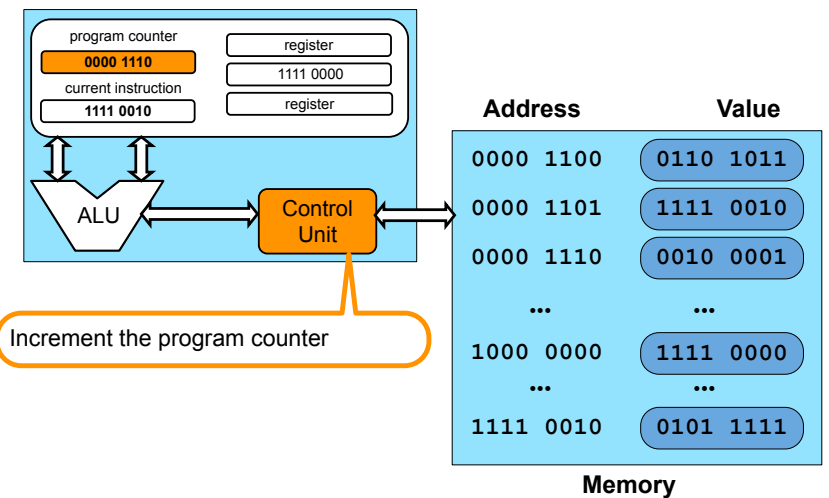
Fetch-Decode-Execute



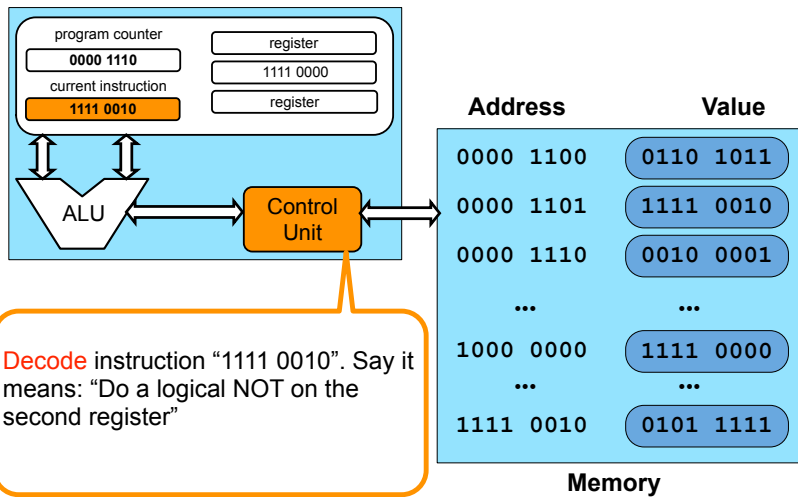
Fetch-Decode-Execute



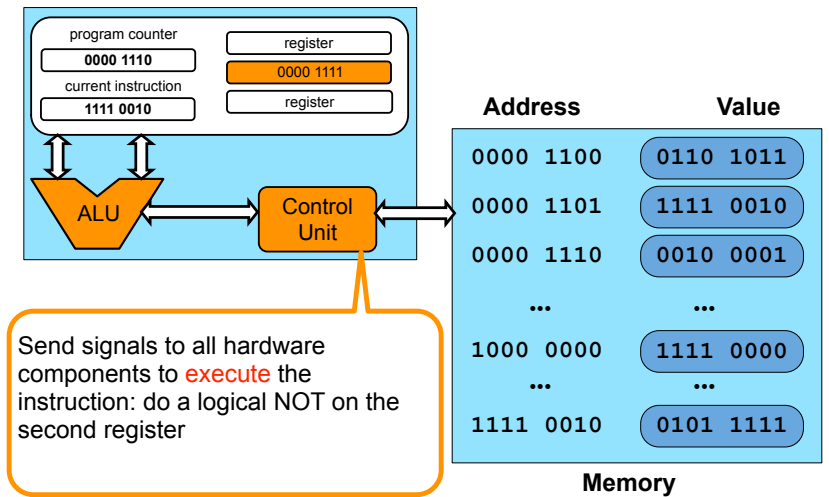
Fetch-Decode-Execute



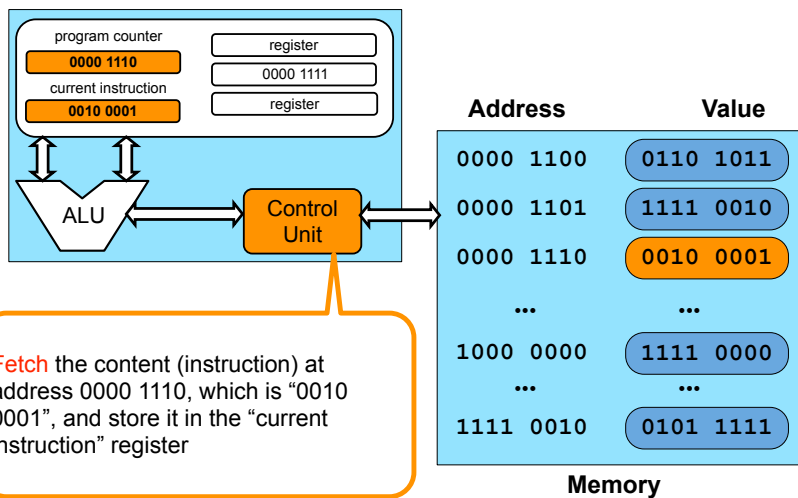
Fetch-Decode-Execute



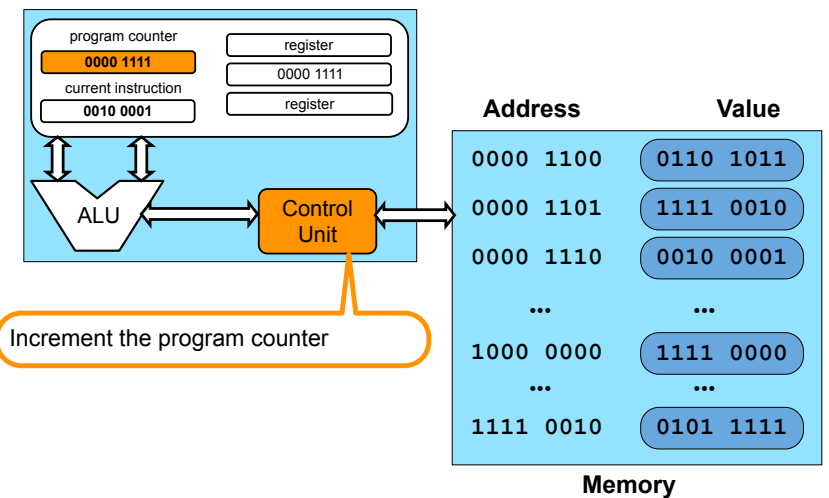
Fetch-Decode-Execute



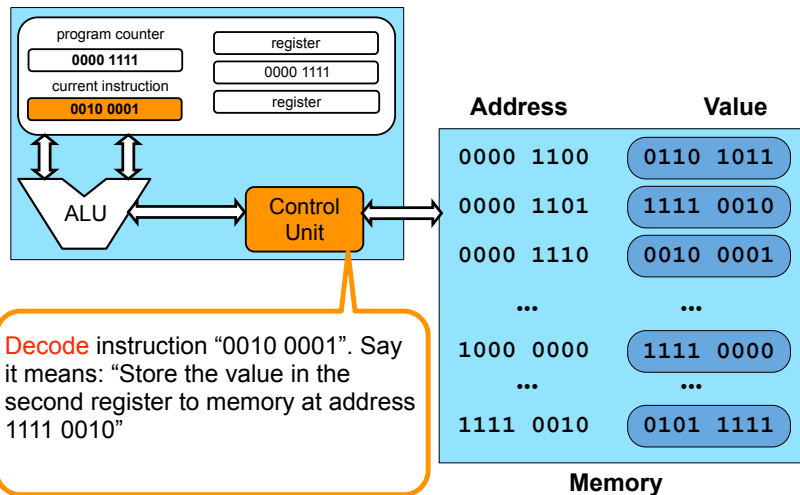
Fetch-Decode-Execute



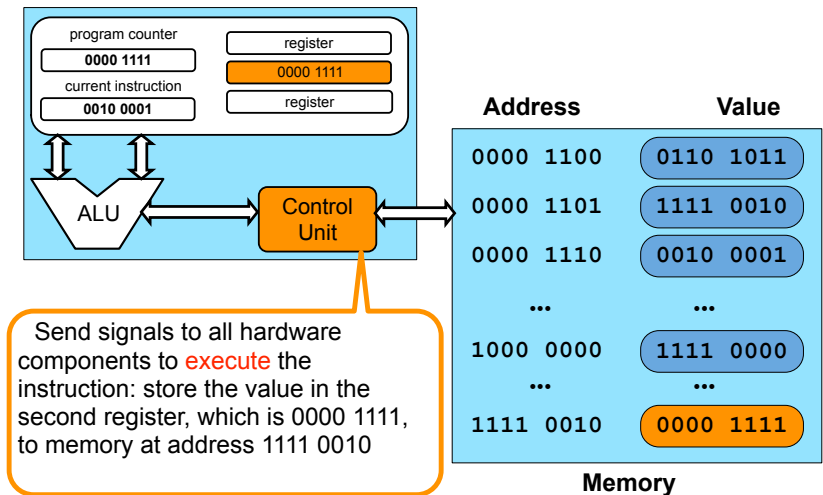
Fetch-Decode-Execute



Fetch-Decode-Execute



Fetch-Decode-Execute




Fetch-Decode-Execute

- This is only a simplified view of the way things work
- The "control unit" is not a single thing
 - Control and data paths are implemented by several complex hardware components
- There are multiple ALUs, there are caches, there are multiple CPUs in fact ("cores")
- Execution is pipelined: e.g., while one instruction is fetched, another is executed
- Decades of computer architecture research have gone into improving performance, thus often leading to staggering hardware complexity
 - Doing smart things in hardware requires more logic gates and wires, thus increasing processor cost
- But conceptually, fetch-decode-execute is it

In-Class Exercise


- With the following instruction set definition and machine state, what is the new memory state after execution completes?

code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address not(register)




- Fetch the instruction: "1111 0000"

code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address not(register)




- Fetch the instruction: "1111 0000"
- Execute it: increment register to value "1111 0011"

code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address not(register)




- Fetch the instruction: "1111 0000"
- Execute it: increment register to value "1111 0011"
- Fetch the next instruction: "1111 0001"

code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address not(register)



- Fetch the instruction: "1111 0000"
- Execute it: increment register to value "1111 0011"
- Fetch the next instruction: "1111 0001"
- Execute it: save value "1111 0011" to address "0000 1100"

code	operation
1111 0000	Increment the register
1111 0010	Decrement the register
0101 1111	Save register to address not(register)



The Clock


- Every computer maintains an internal clock that regulates how quickly instructions can be executed, and is used to synchronize system components
 - Just like a metronome
- In the previous example, each “event” happens at a different “tick” of the clock
- The frequency of the clock is called the **clock rate**
- The time in between two clock ticks is called a clock cycle or **cycle** for short
- Clock cycle = $1 / \text{Clock Rate}$
 - Clock rate = 2.4 GHz
 - Clock cycle = $1 / (2.4 \times 1000 \times 1000 \times 1000)$
= $0.416 \text{ e}^{-9} \text{ sec}$
= 0.416 ns (nanosec)

Faster/slower Clock Rate

- **The higher the clock rate, the shorter the clock cycle**
- It's tempting to think that a faster clock rate means a faster computer
- But it all depends of what amount of work is done in a clock cycle!
 - Computer A: clock rate of 2GHz and a multiplication requires 10 cycles
 - Computer B: clock rate of 1.5GHz and a multiplication requires 5 cycles
 - Computer B is faster than Computer A to run a program that performs a lot of multiplications
- Therefore, clock rates should not be used to compare computers in different families
 - A 1.4GHz Pentium 4 is most likely slower than a 1.5GHz Pentium 4
 - A 2.4GHz Pentium 4 may be slower than a 2.0GHz AMD Athlon64
- Furthermore, comparisons depends on the type of applications
 - Computer A faster than Computer B for some applications
 - Computer B faster than Computer A for some others

Instructions

- Instructions are **encoded** in binary **machine code**
 - e.g.: 01000110101101 may mean “perform an addition of two registers and store the results in another register”
- The CPU is built using **gates** (OR, AND, etc.) which themselves use transistors
 - See ICS331
- These gates implement **instruction decoding**
 - Based on the bits of the instruction code, several signals are sent to different electronic components, which in turn perform useful tasks
- Typically, an instruction consists of two parts
 - The **opcode**: what the instruction computes
 - The **operands**: the input to the computation



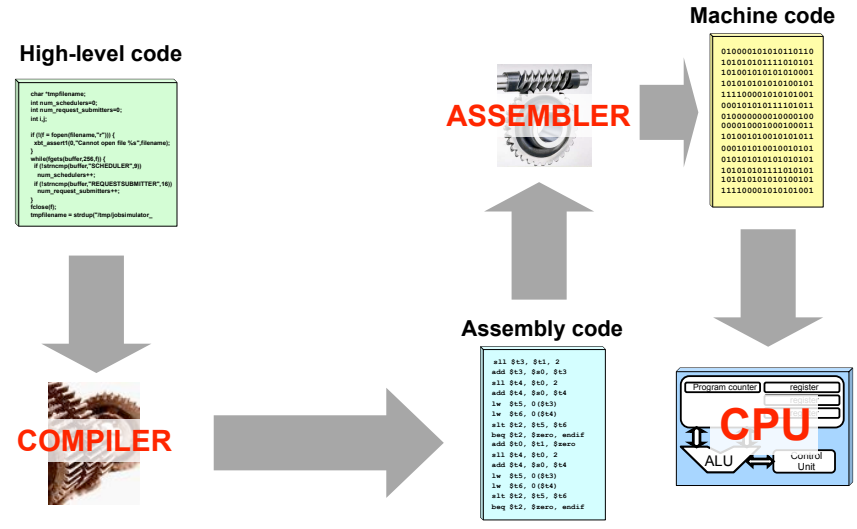
Instruction Set Architecture (ISA)

- When designing a CPU, one must define the set of all the instructions it understands
 - This is one thing that Intel engineers do
- This is called the ISA: Instruction Set Architecture
- Typical ISA include instructions for
 - Performing arithmetic operations on register values
 - Load values from memory into registers
 - Store values from registers into memory
 - Test register values to decide what instruction to execute next
 - ...
- Envision a loooong specification manual that lists all the possible instructions...

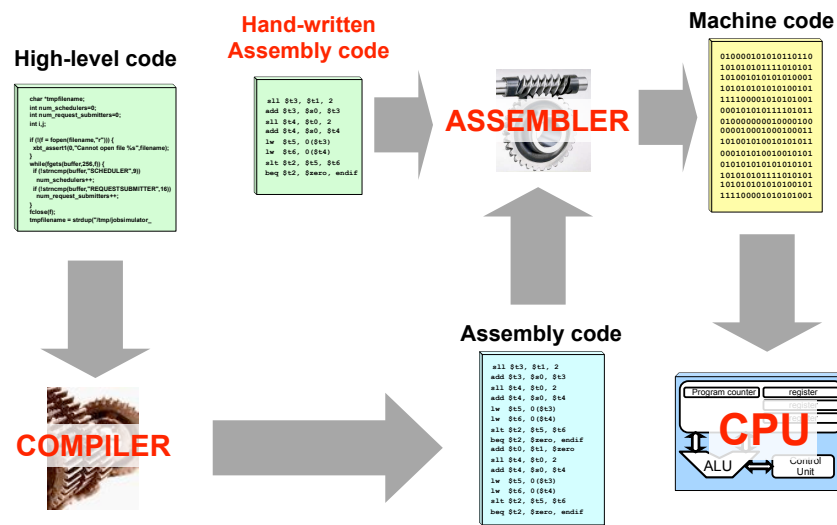
High-Level Languages

- Having high-level languages is good, but CPUs do not understand them
 - As we saw, they only understand very basic instructions to manipulate registers, etc.
- Therefore, there needs to be a **translation** from a high-level language to machine code
- The translation is done by a **compiler**
- Let's see this on a picture....

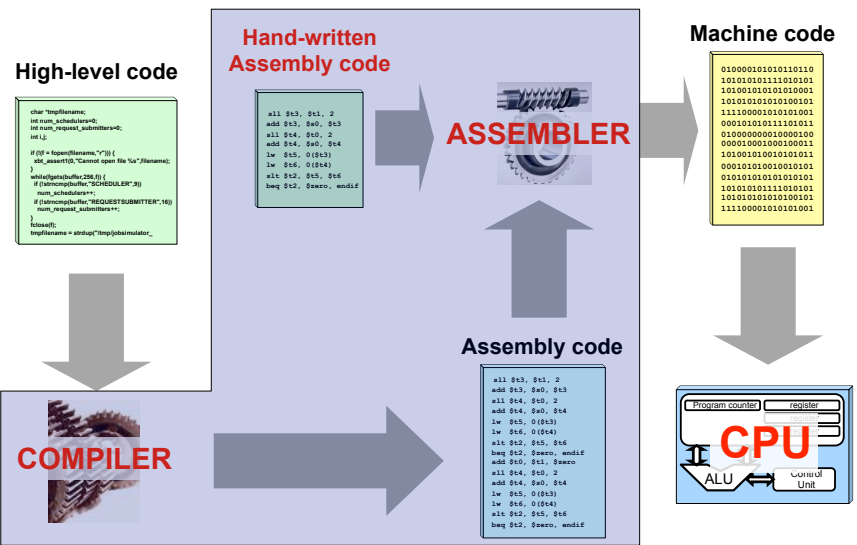
The Big (Simplified) Picture



The Big (Simplified) Picture



This course's topics:



What we do in this course

- **First part of the semester (bulk of the course)**
 - Learn how to write assembly code
 - For the x86 architecture
 - Learn how to use an assembler and a compiler to run our assembly code
- **Second part of the semester (shorter)**
 - Learn about systems tools
 - loader, linker, compiler, debugger, etc.

Why should we learn all this?

- **Why should we learn how to write assembly code?**
 - Students: “We won’t write assembly code for a living!”
- **Reason #1:** Many of you will have to write *some* assembly
 - Write small piece of assembly for performance optimization as part of larger software projects
 - Write assembly code for embedded devices
- **Reason #2:** Learning assembly makes you a better programmer in high-level languages
 - Makes you keenly aware of what happens under the cover, which allows for easier debugging
 - Makes you understand “performance bugs”
 - Allows you to write more efficient high-level code
 - Allows you to read generated assembly to better understand what’s going on

Why should we learn all this?

- **Why should we learn how compilers work?**
 - Students: “We won’t develop compilers for a living!”
- **Reason #1:** Many of you will develop “some” compilers
 - Some of you may develop a compiler for a programming language
 - But often one has to write “compiler” for things that one doesn’t always think of as programming languages
 - E.g., configuration files for large software systems
- **Reason #2:** Knowing how a compiler works makes you a better programmer
 - You know understand the connection between high-level code and generated assembly code (see previous slide)
 - You understand what some high-level language constructs really entail under the cover, and thus understand their performance implications

Why should we learn all this?

- **Meta-reason:** this course should go a long way in giving you a **holistic understanding** of how a program goes from just a text file to a running code
 - You should be able to describe in low-level details how you go from “I wrote a piece of C code that calls a function that adds 2 and 2 together and prints the result” to “the computer prints 4”
 - The complexity of such a simple thing is actually quite stunning, and we’ll take a simplified view
 - There should be something satisfying in knowing how things work from top to bottom!
- This “holistic understanding” should be acquired with ICS312, ICS331, ICS431/EE461, and ICS332

Conclusion

- If you want to know more
 - Take a computer architecture course
 - Classic Textbook: Computer Organization and Design, Fourth Edition: The Hardware/Software Interface (Patterson and Hennessy, Morgan Kaufmann)
- Reading assignment: Sections 1.2 and 1.3
- We'll have a Quiz on this and the previous set of lecture notes
 - this coming Tuesday

